

# Case Study

*Load in the important libraries.*

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr    1.5.1
## v ggplot2    3.5.1      v tibble     3.2.1
## v lubridate  1.9.3      v tidyr      1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(dplyr)
library(tidyr)
library(ggplot2)
library(lubridate)
```

## ***Problem 1 – Data handling, analysis and plotting***

*The first problem of the case study builds on the data in the files `p01-02_portfolio.csv` and `p01-02_rates.csv`. One file contains membership information for a Group Life portfolio and one has information on the rates which should be charged.*

```
# Load the CSV file
portfolio_data <- read_delim(
  "Case Study/data/p01-02_portfolio.csv",
  delim = ";",
  show_col_types = FALSE
)

# View the data
head(portfolio_data)
```

```
## # A tibble: 6 x 5
##   SchemeName Date.of.Birth Gender DeathSI Industry
##   <chr>      <chr>      <chr> <chr>   <chr>
## 1 Scheme2    29.05.1949   F     <NA>    Government & Public Administration
## 2 Scheme2    07.09.1950   F     <NA>    Government & Public Administration
## 3 Scheme2    27.09.1956   F     <NA>    Government & Public Administration
## 4 Scheme2    18.02.1942   F     <NA>    Government & Public Administration
## 5 Scheme2    31.07.1951   F     <NA>    Government & Public Administration
## 6 Scheme2    10.07.1960   F     <NA>    Government & Public Administration
```

```
# Load the CSV file
rates_data <- read_delim("Case Study/data/p01-02_rates.csv",
                        delim = ";",
                        show_col_types = FALSE)

# View the data
head(rates_data)
```

```
## # A tibble: 6 x 3
##   Age Gender Rate
##   <dbl> <chr> <dbl>
## 1    18 M     0.32
## 2    19 M     0.32
## 3    20 M     0.32
## 4    21 M     0.31
## 5    22 M     0.31
## 6    23 M     0.31
```

*rates\_data*

```
# Count occurrences of each combination of Gender and Age
duplicates <- rates_data %>%
  group_by(Gender, Age) %>%
  summarise(count = n()) %>%
  filter(count > 1)
```

```
## 'summarise()' has grouped output by 'Gender'. You can override using the
## '.groups' argument.
```

```
# Check if any duplicates exist
if (nrow(duplicates) == 0) {
  message("Sanity Check Passed: 'Gender' and 'Age' form a unique key.")
} else {
  message("Sanity Check Failed: There are duplicate combinations of 'Gender' and 'Age'.")
  print(duplicates)
}
```

```
## Sanity Check Passed: 'Gender' and 'Age' form a unique key.
```

### Question a.

Read the data from the two files into R's memory. The rates are applicable to each individual in the portfolio, depending on that individual's age and gender. Combine the two datasets into a single table by looking up the rate for each line of the portfolio.

```
# Step 1: Convert the Date.of.Birth column to Date format
portfolio_data$Date.of.Birth <- dmy(portfolio_data$Date.of.Birth) # dmy is used for "day-month-year" f

# Step 2: Calculate the time difference in years
portfolio_data$age <- ceiling(interval(portfolio_data$Date.of.Birth, today()) / years(1))

# Step 3: View the updated data with age column
head(portfolio_data)
```

```
## # A tibble: 6 x 6
##   SchemeName Date.of.Birth Gender DeathSI Industry age
##   <chr>      <date>      <chr> <chr>   <chr>      <dbl>
## 1 Scheme2    1949-05-29    F    <NA>    Government & Public Administrat~ 76
## 2 Scheme2    1950-09-07    F    <NA>    Government & Public Administrat~ 75
## 3 Scheme2    1956-09-27    F    <NA>    Government & Public Administrat~ 68
## 4 Scheme2    1942-02-18    F    <NA>    Government & Public Administrat~ 83
## 5 Scheme2    1951-07-31    F    <NA>    Government & Public Administrat~ 74
## 6 Scheme2    1960-07-10    F    <NA>    Government & Public Administrat~ 65
```

```
# Inner join the two datasets
combined_data <- inner_join(portfolio_data,
                             rates_data,
                             by = c("age" = "Age", "Gender" = "Gender"),
)

# Check if the row count of the joined data matches the original portfolio data
if (nrow(combined_data) == nrow(portfolio_data)) {
  message("Sanity Check Passed: The row count of combined_data matches portfolio_data.")
} else {
  message("Sanity Check Failed: The row count of combined_data does not match portfolio_data.")
  message("Rows in portfolio_data: ", nrow(portfolio_data))
  message("Rows in combined_data: ", nrow(combined_data))
}
```

```
## Sanity Check Failed: The row count of combined_data does not match portfolio_data.
```

```
## Rows in portfolio_data: 177922
```

```
## Rows in combined_data: 145607
```

*This sanity check is expected to fail because the rates data is cutoff at 70. Do not consider people over 70 in this analysis.*

```
# Find rows in portfolio_data that do not have a match in rates_data
missing_matches <- anti_join(portfolio_data,
                              rates_data,
                              by = c("age" = "Age", "Gender" = "Gender"))

# Check how many rows are missing
nrow(missing_matches)
```

*Investigate missing matches*

```
## [1] 32315
```

### Question b.

Group the Industry field into common-sense based groupings and determine the mean, standard deviation and quantiles of DeathSI for each of your industry groups.

```
industry_counts <- combined_data %>%
  count(Industry) %>%
  arrange(desc(n))

# View the result
print(industry_counts)
```

```
## # A tibble: 33 x 2
##   Industry          n
##   <chr>          <int>
## 1 <NA>          87405
## 2 Government & Public Administration 29475
## 3 Other         14130
## 4 Sporting Club   2177
## 5 Ex-Services Club 1442
## 6 BSS-Business Services 1179
## 7 MAN-Manufacturing 1065
## 8 EDN-Education    874
## 9 COM-Communication Serv. 817
## 10 FIN-Finance & Insurance 779
## # i 23 more rows
```

```
combined_data <- combined_data %>%
  mutate(
    Industry_Group = case_when(
      Industry %in% c("Government & Public Administration", "Ex-Services Club") ~ "Government and Public Administration",
      Industry %in% c(
        "Sporting Club",
        "Golf Club",
        "Bowls Club",
        "Registered Club",
        "Surf Life Saving Club",
        "Workers Club",
        "Australian Rules Football Club",
        "Leagues Club",
        "Associated with Club Industry"
      ) ~ "Clubs and Associations",
      Industry %in% c(
        "BSS-Business Services",
        "FIN-Finance & Insurance",
        "Professional Services",
        "LAW-Solicitors/Barrister",
        "ENG-Engineers",
        "MGE-Medical Services Gen"
      ) ~ "Professional and Business Services",
      Industry %in% c(
        "MAN-Manufacturing",
        "CON-Construction",
```

```

      "ELE-Electricians",
      "VEH-Vehicle Industry",
      "WEO-Wholesale Trades"
    ) ~ "Manufacturing, Construction, and Trades",
    Industry %in% c(
      "EDN-Education",
      "HEA-Health Industry",
      "MGE-Medical Services Gen"
    ) ~ "Education and Health",
    Industry %in% c(
      "RTL-Retail Trade",
      "ACR-Accom. Cafes & Rests",
      "FOO-Food",
      "Hospitality"
    ) ~ "Retail, Hospitality, and Food",
    Industry %in% c("AGR-Farming/Agriculture", "EGW
-Electric/Gas/Water") ~ "Agriculture and Utilities",
    Industry == "Other" ~ "Other",
    TRUE ~ "Uncategorized" # Catch any uncategorized industries
  )
)

# View the newly grouped data
print(combined_data)

```

```

## # A tibble: 145,607 x 8
##   SchemeName Date.of.Birth Gender DeathSI Industry    age  Rate Industry_Group
##   <chr>      <date>      <chr>  <chr>  <chr>    <dbl> <dbl> <chr>
## 1 Scheme2    1956-09-27    F    <NA>  Governmen~ 68  5.96 Government an~
## 2 Scheme2    1960-07-10    F    <NA>  Governmen~ 65  4.35 Government an~
## 3 Scheme2    1954-12-24    F    <NA>  Governmen~ 70  7.34 Government an~
## 4 Scheme2    1958-02-28    F    <NA>  Governmen~ 67  5.36 Government an~
## 5 Scheme2    1968-09-12    F    <NA>  Governmen~ 57  1.91 Government an~
## 6 Scheme2    1966-11-21    F    <NA>  Governmen~ 58  2.11 Government an~
## 7 Scheme2    1957-03-05    F    <NA>  Governmen~ 68  5.96 Government an~
## 8 Scheme2    1966-02-01    F    <NA>  Governmen~ 59  2.34 Government an~
## 9 Scheme2    1975-02-10    F    <NA>  Governmen~ 50  0.96 Government an~
## 10 Scheme2   1966-11-21    F    <NA>  Governmen~ 58  2.11 Government an~
## # i 145,597 more rows

```

```

industry_counts <- combined_data %>%
  count(Industry_Group) %>%
  arrange(desc(n))

# View the result
print(industry_counts)

```

```

## # A tibble: 9 x 2
##   Industry_Group      n
##   <chr>          <int>
## 1 Uncategorized    88467
## 2 Government and Public Services 30917

```

```
## 3 Other 14130
## 4 Clubs and Associations 4805
## 5 Manufacturing, Construction, and Trades 2376
## 6 Professional and Business Services 2364
## 7 Education and Health 1242
## 8 Retail, Hospitality, and Food 974
## 9 Agriculture and Utilities 332
```

```
1. # Check the type of DeathSI
   typeof(combined_data$DeathSI)
```

```
## [1] "character"
```

```
# Count the number of NA values in DeathSI when it was character type
na_count <- sum(is.na(combined_data$DeathSI))
na_count
```

```
## [1] 14204
```

```
# Count the number of "NA" string values in DeathSI when it was character type
na_string_count <- sum(combined_data$DeathSI == "NA", na.rm = TRUE)
na_string_count
```

```
## [1] 0
```

```
# Remove apostrophes and convert the DeathSI column from character to numeric
combined_data$DeathSI <- as.numeric(gsub("'", "", combined_data$DeathSI))
```

```
# Check the type of DeathSI
typeof(combined_data$DeathSI)
```

```
## [1] "double"
```

```
# Count the number of NA values in DeathSI when it is the double type
na_count <- sum(is.na(combined_data$DeathSI))
na_count
```

```
## [1] 14204
```

```
# Calculate mean, standard deviation, and quantiles for each industry group
summary_stats <- combined_data %>%
  group_by(Industry_Group) %>%
  summarize(
    mean_value = mean(DeathSI, na.rm = TRUE),
    sd_value = sd(DeathSI, na.rm = TRUE),
    q25 = quantile(DeathSI, 0.25, na.rm = TRUE),
    median_value = median(DeathSI, na.rm = TRUE),
    q75 = quantile(DeathSI, 0.75, na.rm = TRUE)
  )

# View the result
print(summary_stats)
```

```
## # A tibble: 9 x 6
##   Industry_Group      mean_value sd_value    q25 median_value    q75
##   <chr>              <dbl>    <dbl>  <dbl>      <dbl>    <dbl>
## 1 Agriculture and Utilities    174894.  224538.  3.39e4    142202  2.01e5
## 2 Clubs and Associations      279598.  184125.  2.00e5    220613  3.02e5
## 3 Education and Health       350616.  234855.  1.79e5    320172  4.88e5
## 4 Government and Public Services 224447.   98490.  1.5 e5    220000  3 e5
## 5 Manufacturing, Construction, a~ 303140.  236342.  1.35e5    263191  4.07e5
## 6 Other                    262005.  114205.  2.08e5    245716  2.86e5
## 7 Professional and Business Serv~ 449616.  323093.  2.52e5    383741  5.52e5
## 8 Retail, Hospitality, and Food   320963.  205977.  2.05e5    248697  3.89e5
## 9 Uncategorized              228127.  214916.  9.07e4    170742  2.85e5
```